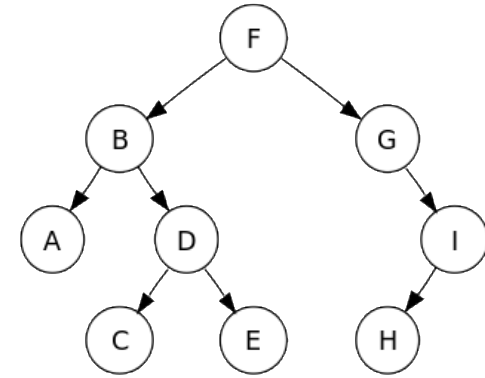


Advanced Data Structure and Algorithm

Binary Search Trees

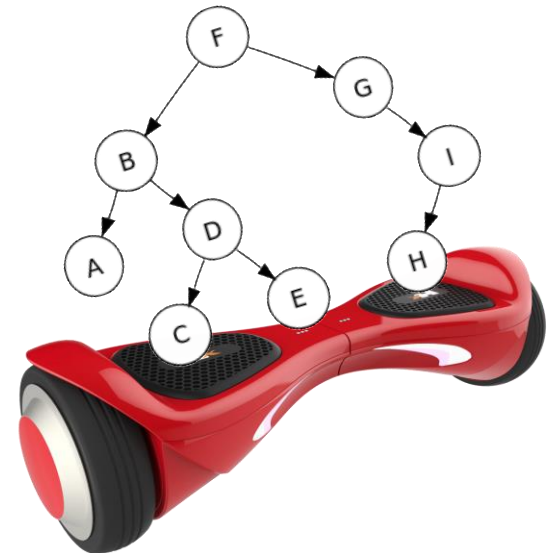
Today

- Binary search trees
 - They are better when they're balanced.



this will lead us to...

- Self-Balancing Binary Search Trees
 - **Red-Black** trees.




Some data structures
for storing objects like **5** (aka, **nodes** with **keys**)

- (Sorted) arrays:



- (Sorted) linked lists:



- Some basic operations:
 - **INSERT, DELETE, SEARCH**

Sorted Arrays



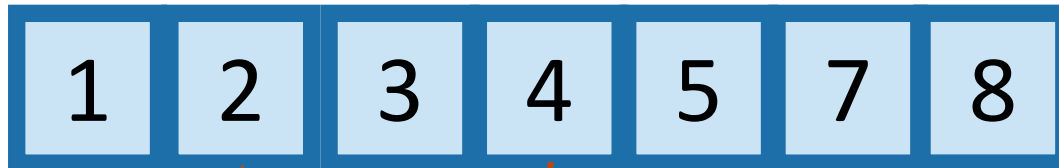
- $O(n)$ INSERT/DELETE:

- First, find the relevant element (time $O(\log(n))$ as below), and then move a bunch elements in the array:



eg, insert 4.5

- $O(\log(n))$ SEARCH:



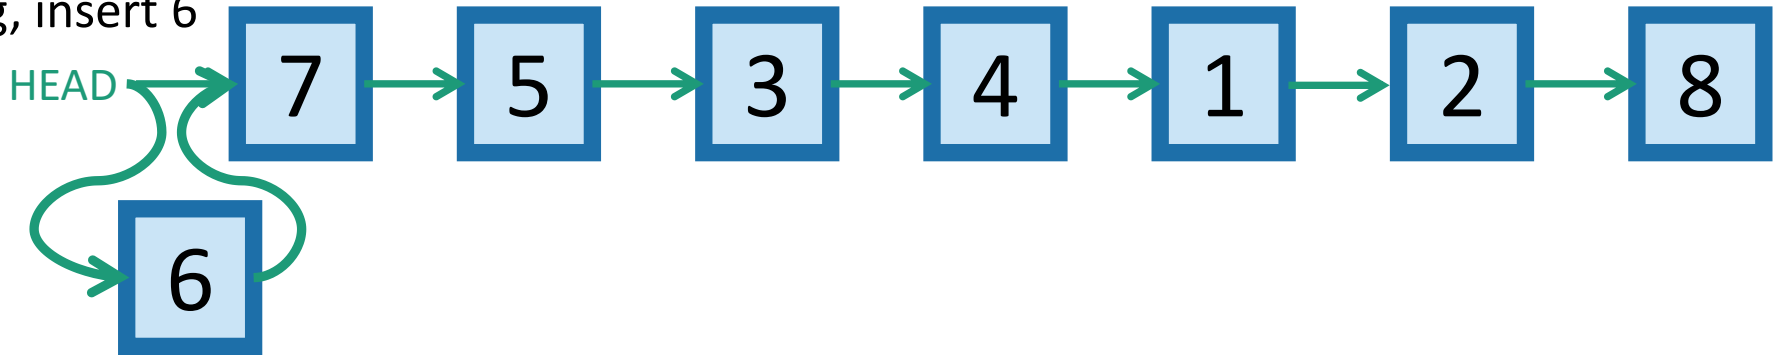
eg, Binary search to see if 3 is in A.

UNSorted linked lists

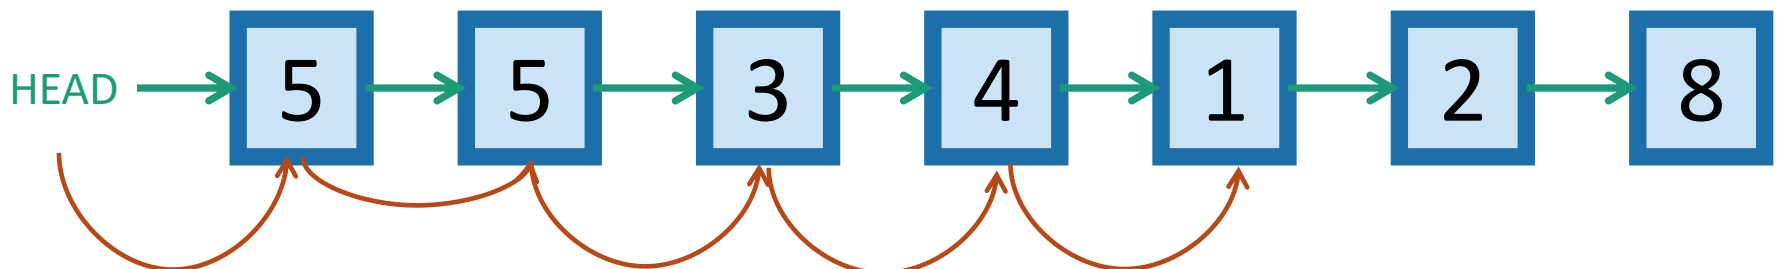


- $O(1)$ INSERT:

eg, insert 6



- $O(n)$ SEARCH/DELETE:



eg, search for 1 (and then you could delete it by manipulating pointers).

Motivation for Binary Search Trees

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$	$O(n)$	$O(\log(n))$
Delete	$O(n)$	$O(n)$	$O(\log(n))$
Insert	$O(n)$	$O(1)$	$O(\log(n))$

Motivation for Binary Search Trees

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 😊	$O(n)$	$O(\log(n))$
Delete	$O(n)$	$O(n)$	$O(\log(n))$
Insert	$O(n)$	$O(1)$ 😊	$O(\log(n))$

Motivation for Binary Search Trees

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 😊	$O(n)$ 😞	$O(\log(n))$
Delete	$O(n)$ 😞	$O(n)$ 😞	$O(\log(n))$
Insert	$O(n)$ 😞	$O(1)$ 😊	$O(\log(n))$

Motivation for Binary Search Trees

TODAY!

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 😊	$O(n)$ 😞	$O(\log(n))$ 😊
Delete	$O(n)$ 😞	$O(n)$ 😞	$O(\log(n))$ 😊
Insert	$O(n)$ 😞	$O(1)$ 😊	$O(\log(n))$ 😊

For today all keys are distinct.

Binary tree terminology

Each node has two **children**.

The **left child** of 3 is 2

The **right child** of 3 is 4

The **parent** of 3 is 5

2 is a **descendant** of 5


Each node has a pointer to its left child, right child, and parent.

Both **children** of 1 are NIL.
(Not usually drawn).

The **height** of this tree is 3.
(Max number of edges from the root to a leaf).

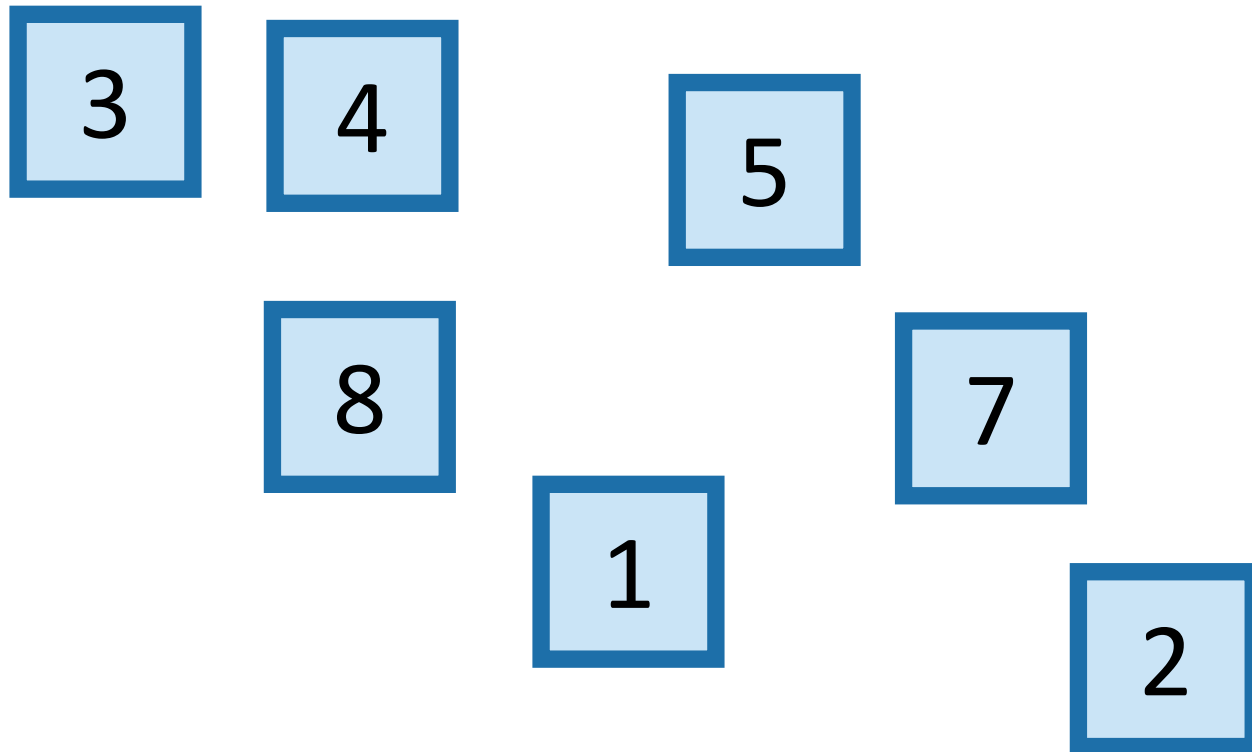
This node is the **root**

This is a **node**.
It has a **key** (7).



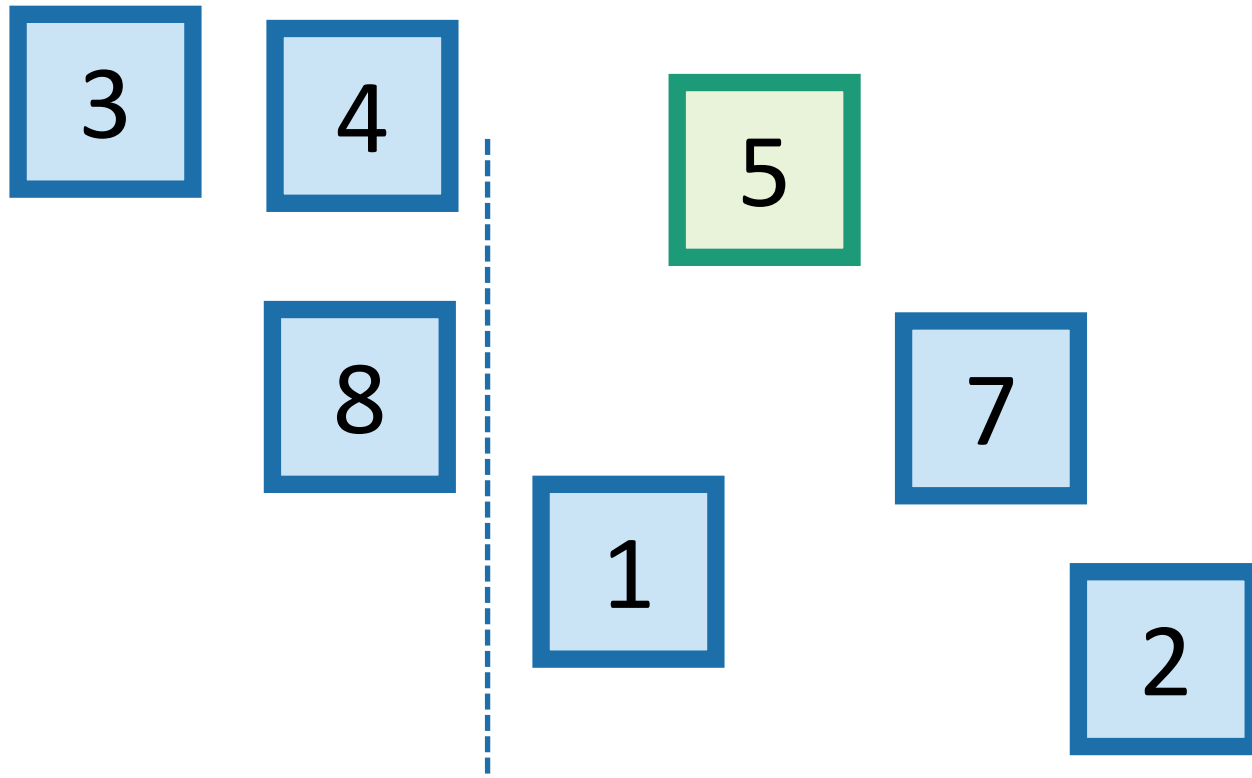
Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



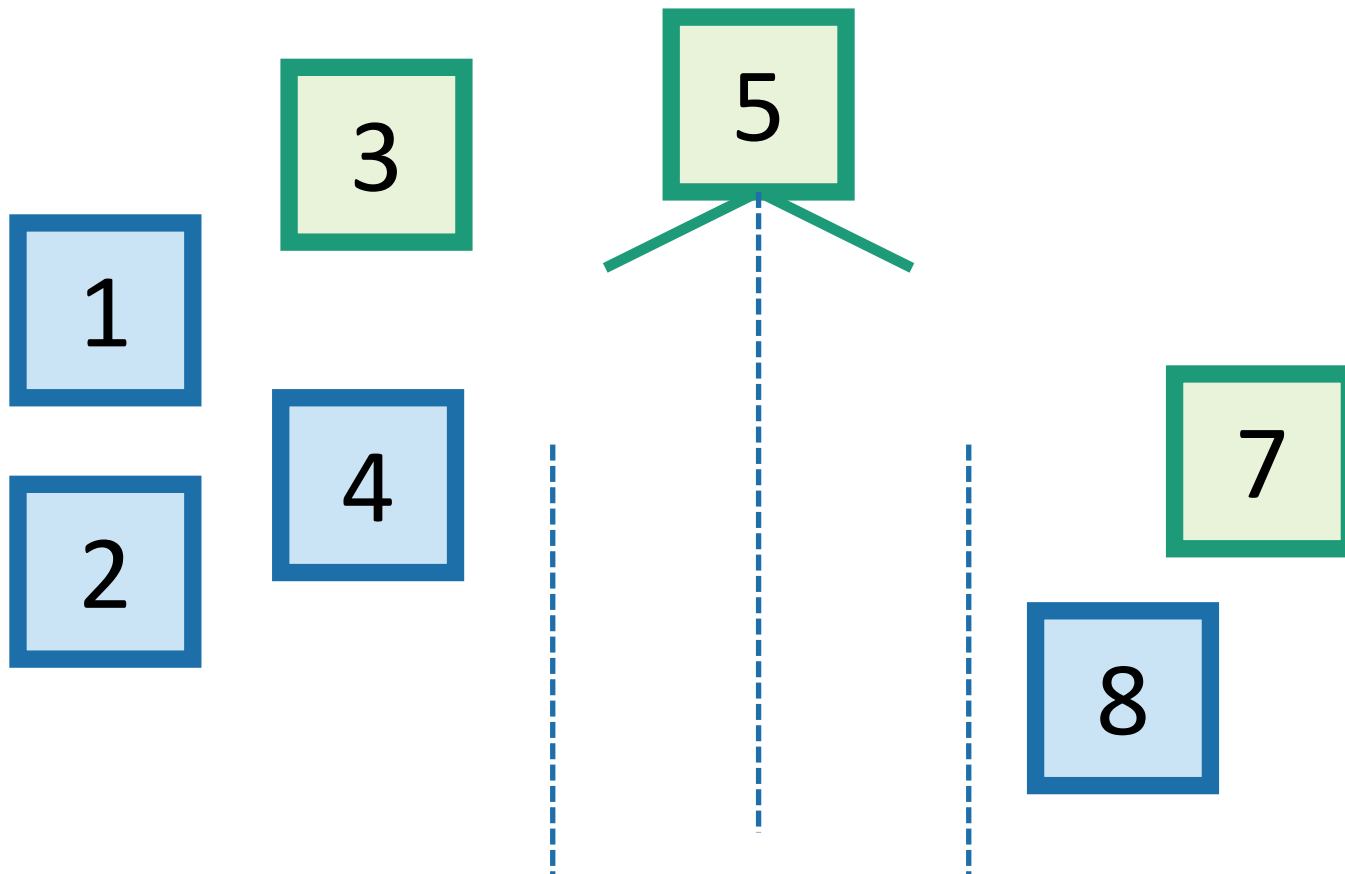
Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



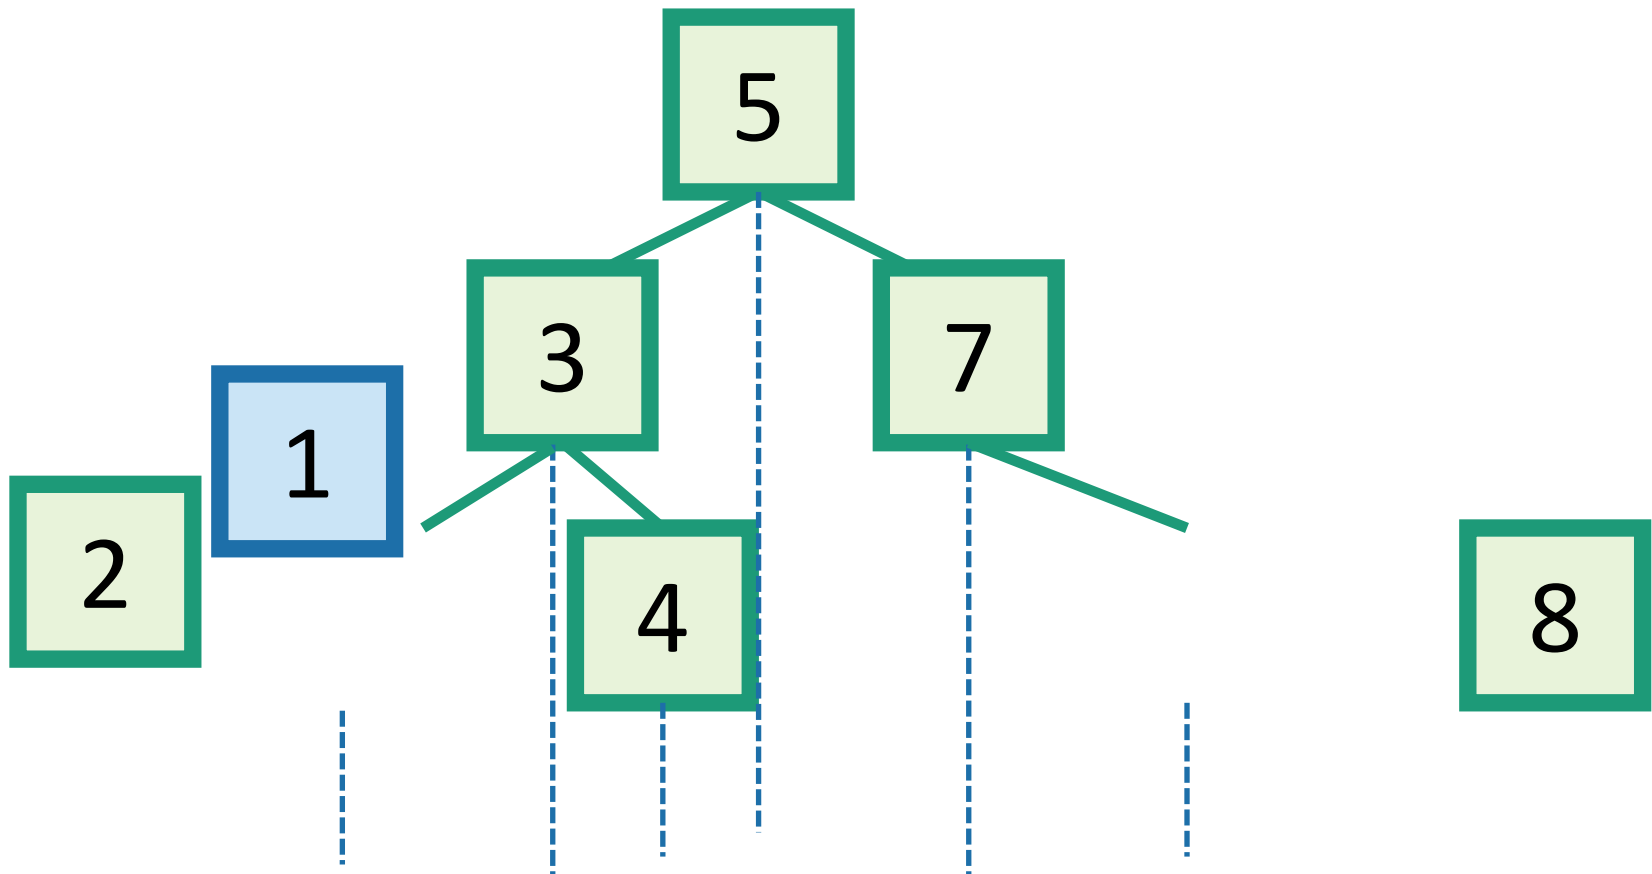
Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



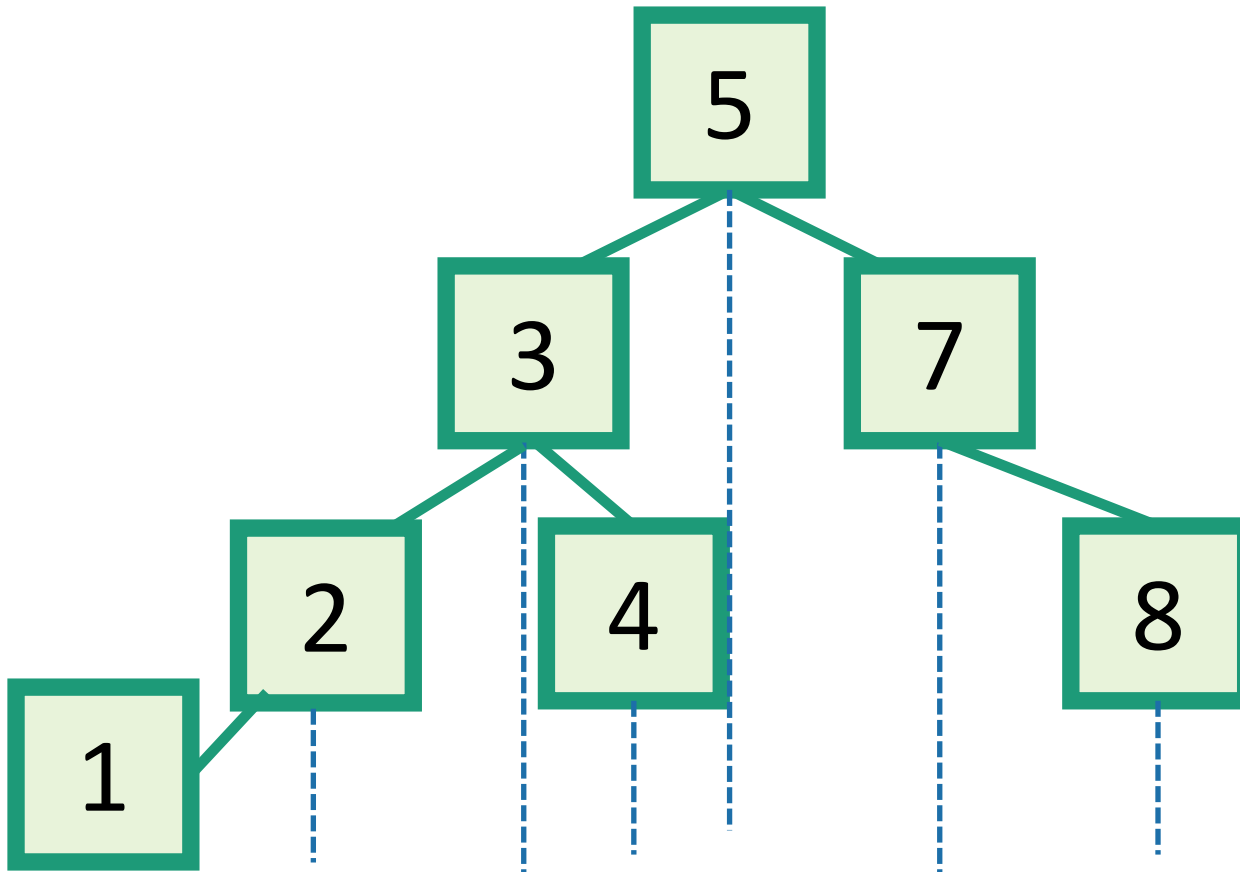
Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:

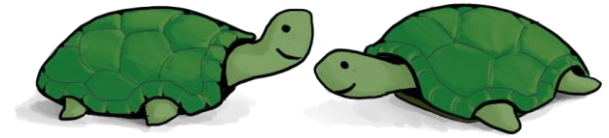


Q: Is this the only binary search tree I could possibly build with these values?



A: **No.** I made choices about which nodes to choose when. Any choices would have been fine.

Binary Search Trees

Which of these is a BST?



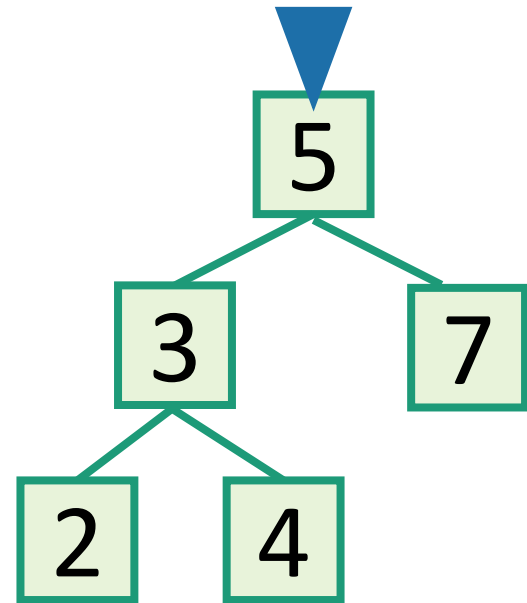
- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.



Aside: In-Order Traversal of BSTs

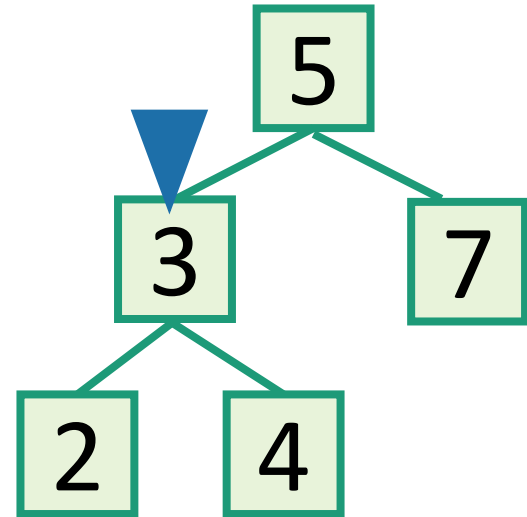
- Output all the elements in sorted order!

- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



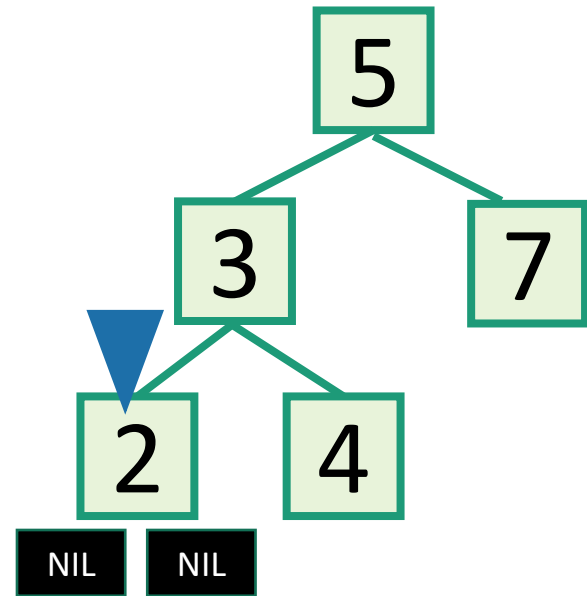
Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



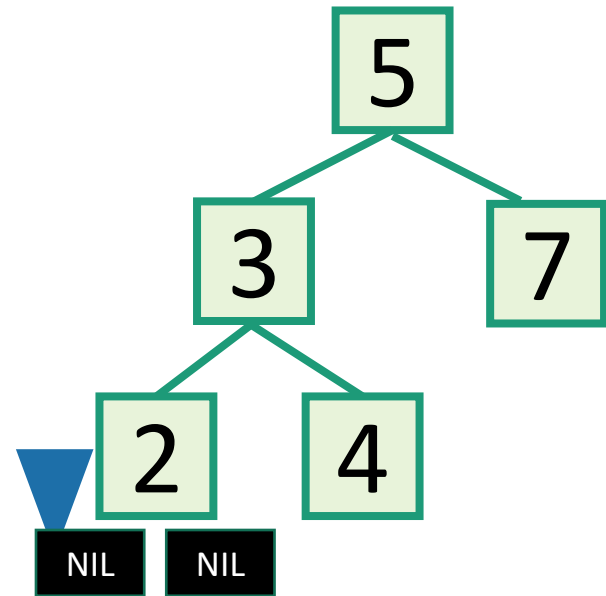
Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



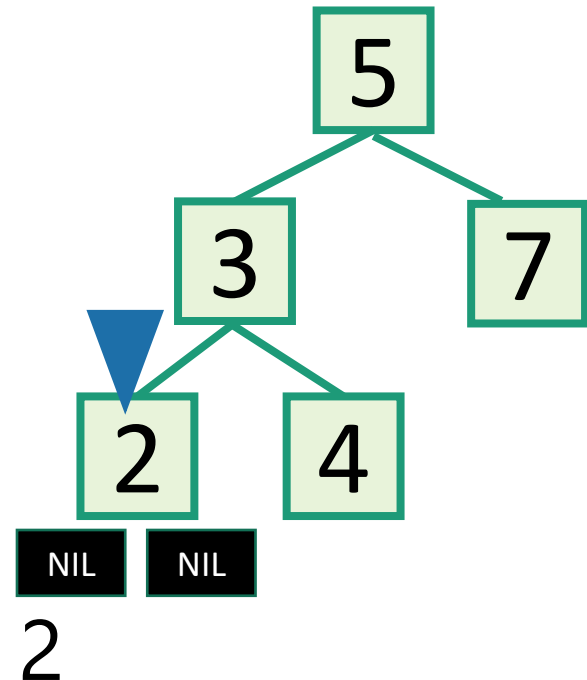
Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



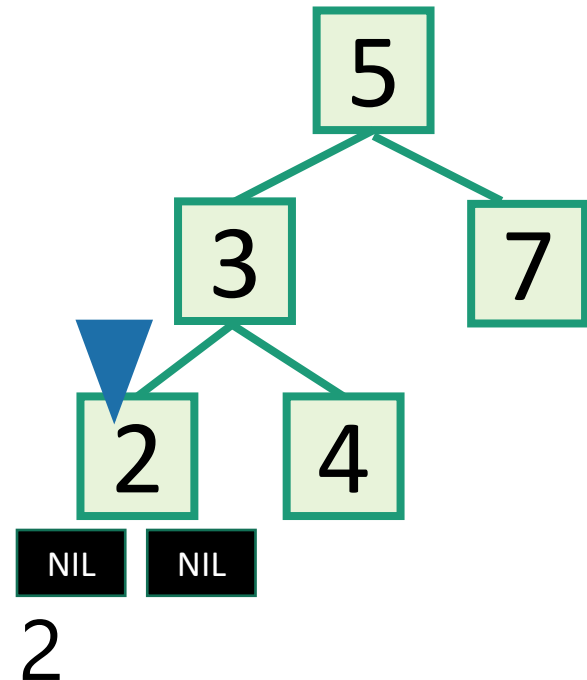
Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!

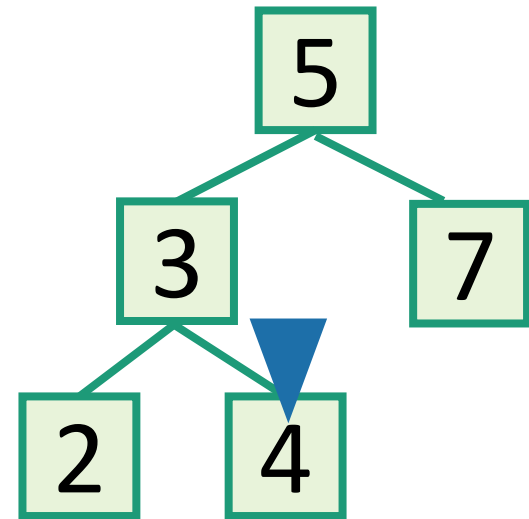
- inOrderTraversal(x):
 - if $x \neq \text{NIL}$:
 - inOrderTraversal(x.left)
 - print(x.key)
 - inOrderTraversal(x.right)



2 3

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



2 3 4

Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`

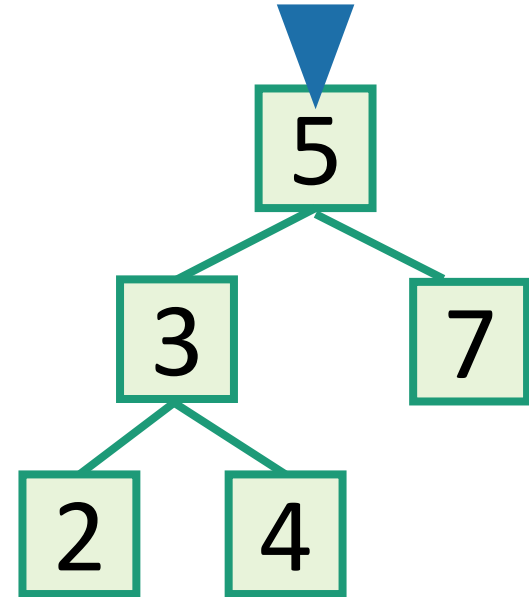


2 3 4

Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!

- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`

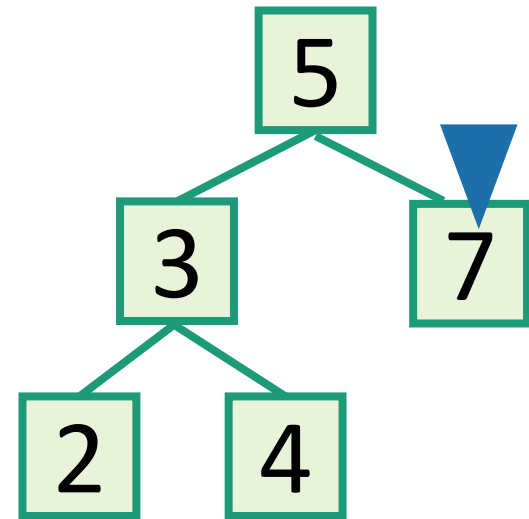


2 3 4 5

Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!

- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`

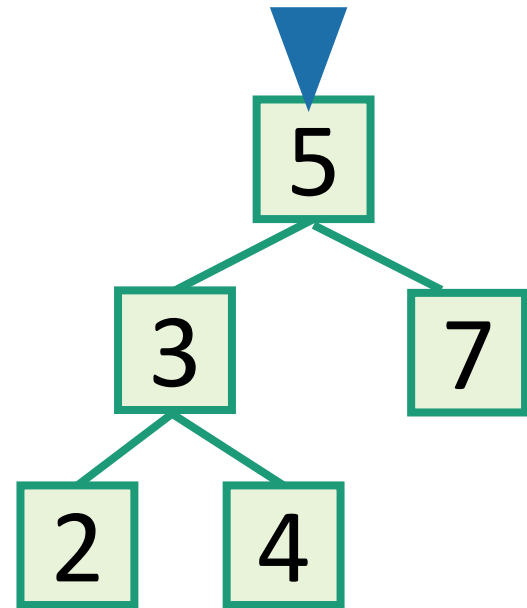


2 3 4 5 7

Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!

- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.key)`
 - `inOrderTraversal(x.right)`



- Runs in time $O(n)$.

2 3 4 5 7 Sorted!

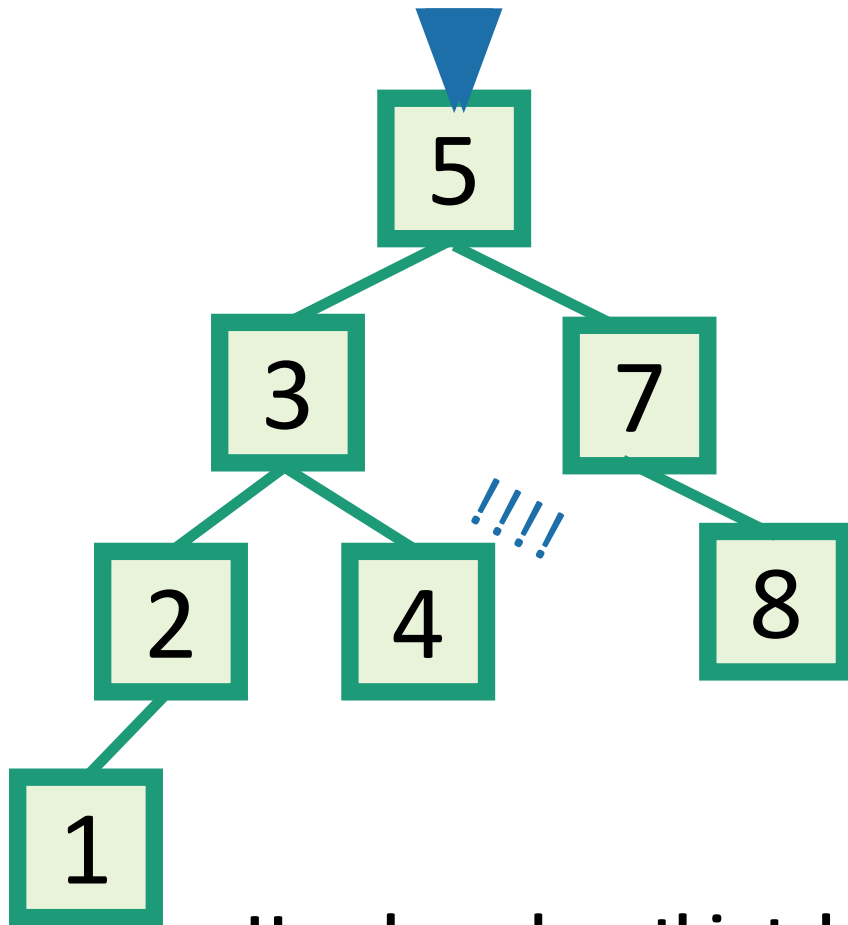
Back to the goal

Fast **SEARCH**/**INSERT**/**DELETE**

Can we do these?

SEARCH in a Binary Search Tree

definition by example



How long does this take?

$O(\text{length of longest path}) = O(\text{height})$

EXAMPLE: Search for 4.


EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

Write pseudocode
(or actual code) to
implement this!



INSERT in a Binary Search Tree

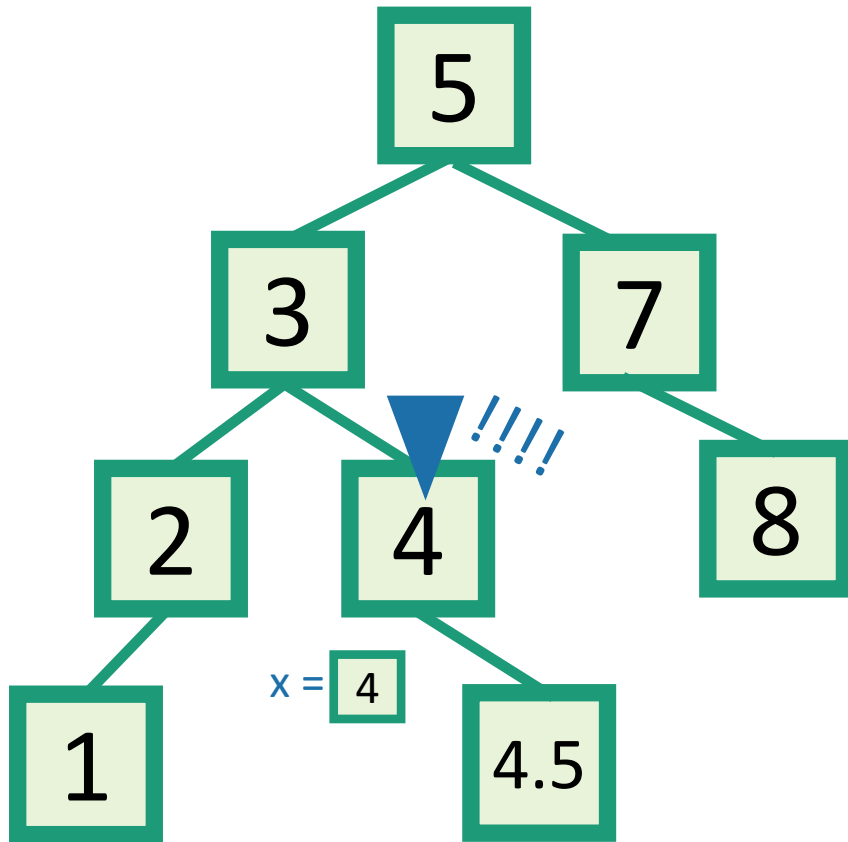


EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **Insert** a new node with desired key at x ...

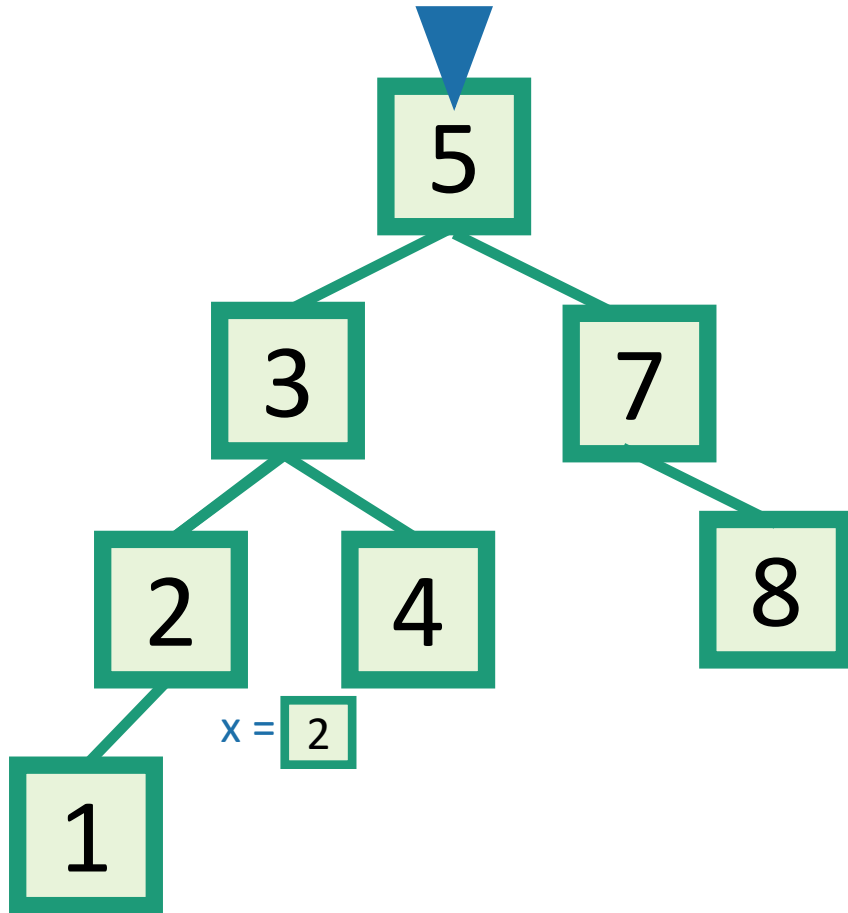
INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5



- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $\text{key} > x.\text{key}$:
 - Make a new node with the correct key, and put it as the right child of x .
 - **if** $\text{key} < x.\text{key}$:
 - Make a new node with the correct key, and put it as the left child of x .
 - **if** $x.\text{key} == \text{key}$:
 - **return**

DELETE in a Binary Search Tree



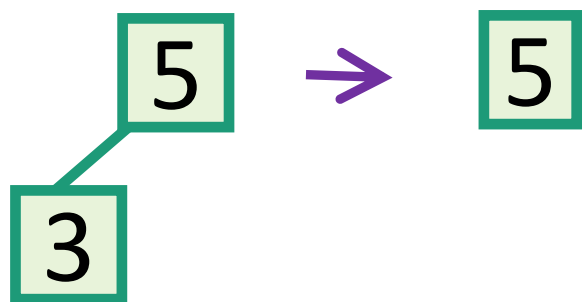
EXAMPLE: Delete 2

- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x

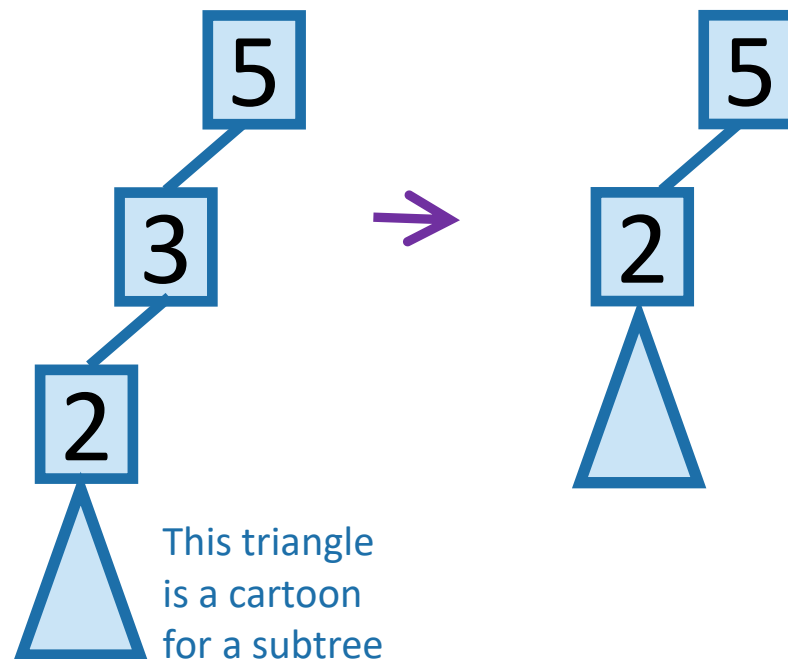
DELETE in a Binary Search Tree

several cases (by example)

say we want to delete 3



Case 1: if 3 is a leaf,
just delete it.




Write pseudocode for all of these!



Case 2: if 3 has just one child,
move that up.

DELETE in a Binary Search Tree

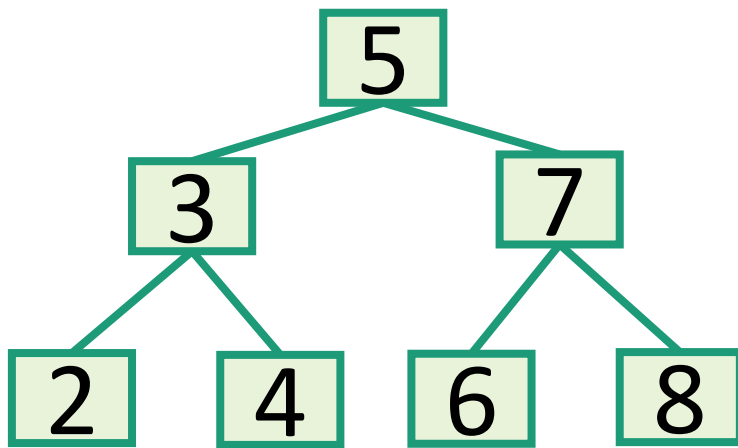
Case 3: if 3 has two children,
replace 3 with its **immediate successor**.
(aka, next biggest thing after 3)



- Does this maintain the BST property?
 - Yes.
- How do we find the immediate successor?
 - SEARCH for 3 in the subtree under 3.right
- How do we remove it when we find it?
 - If [3.1] has 0 or 1 children, do one of the previous cases.
- What if [3.1] has two children?
 - It doesn't. (can not have two children)

How long do these operations take?

- **SEARCH** is the big one.
 - Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.

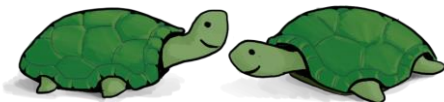


Time = $O(\text{height of tree})$

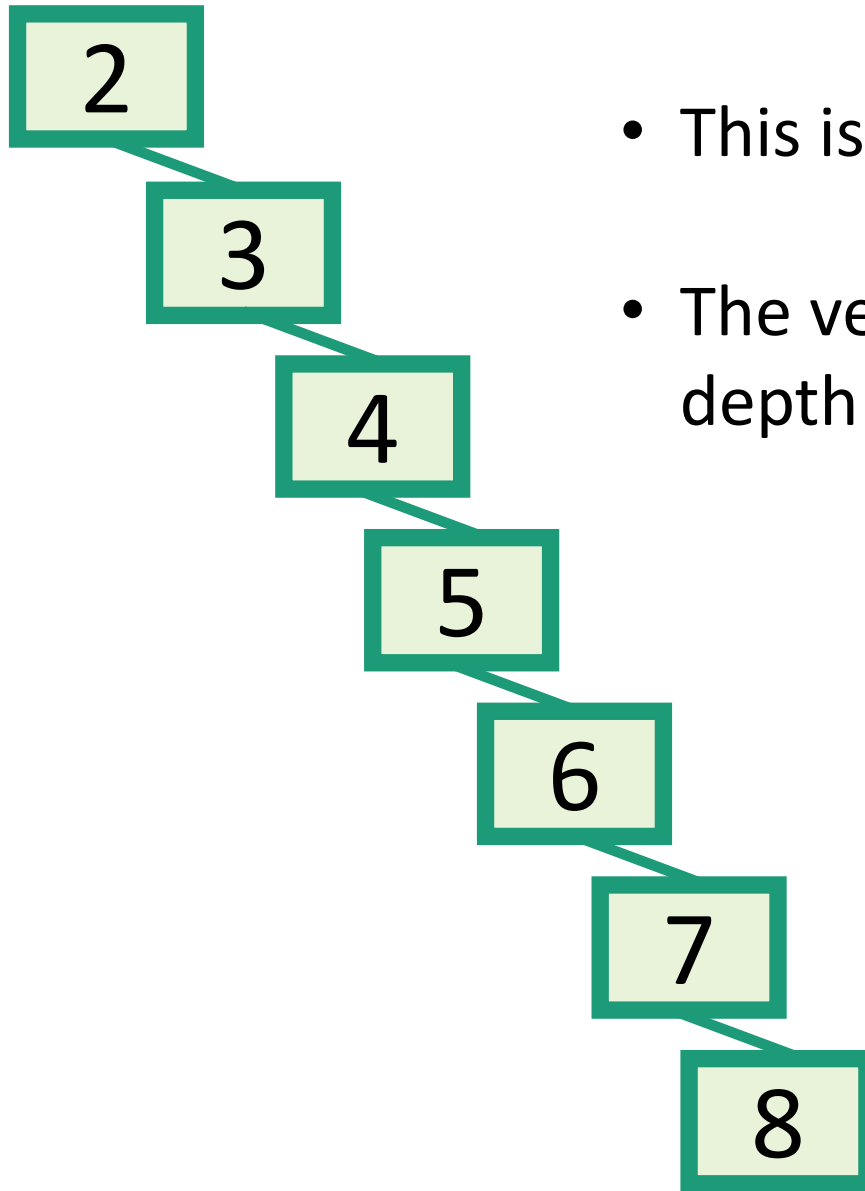
Trees have depth $O(\log(n))$. **Done!**

Wait a second...

How long does search take?



Search might take time $O(n)$.



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.

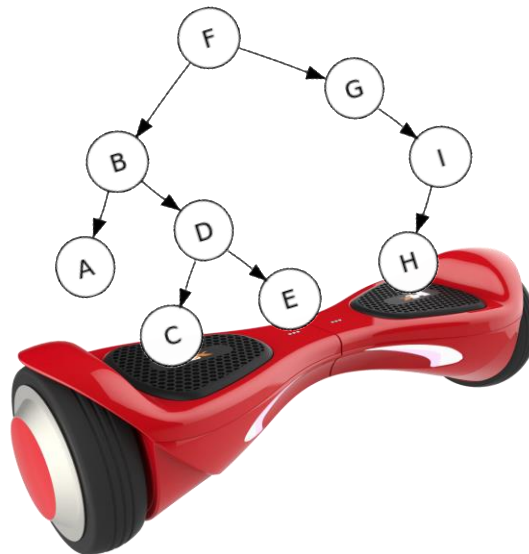
What to do?

How often is “every so often” in the worst case?
It’s actually pretty often!



- Goal: Fast **SEARCH**/**INSERT**/**DELETE**
- All these things take time $O(\text{height})$
- And the height might be big!!! ☹️
- Idea 0:
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
 - At least $\Omega(n)$ every so often....
- Turns out that’s not a great idea. Instead we turn to...

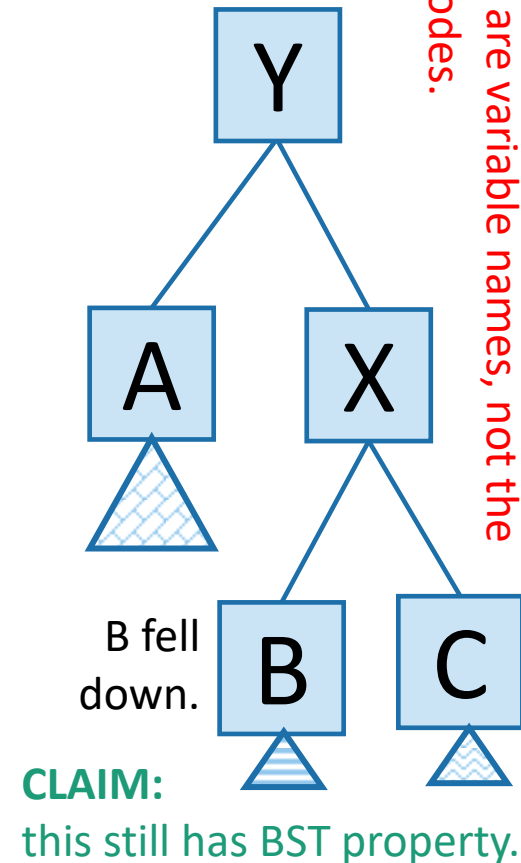
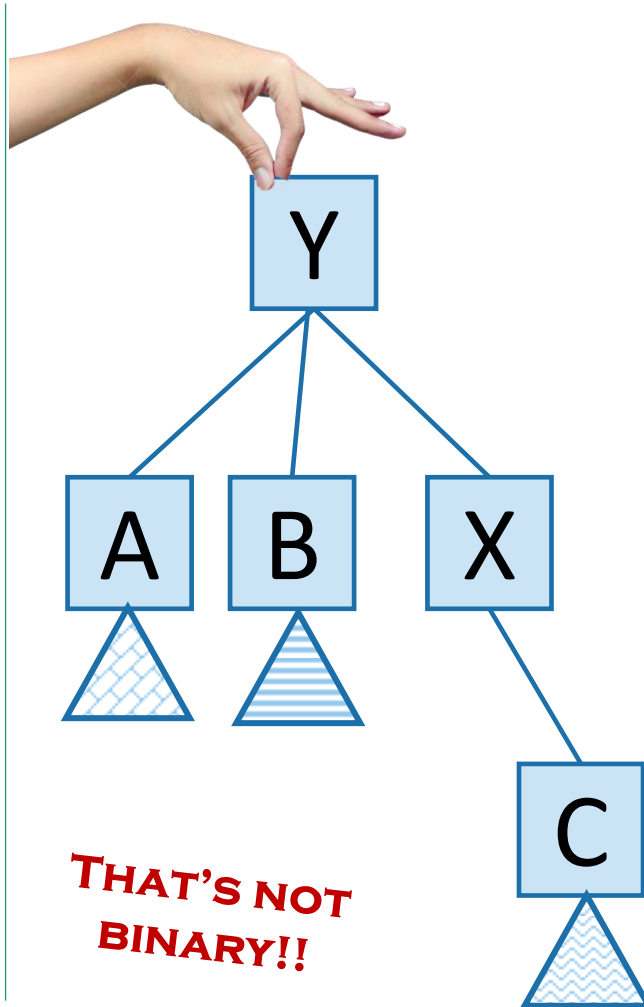
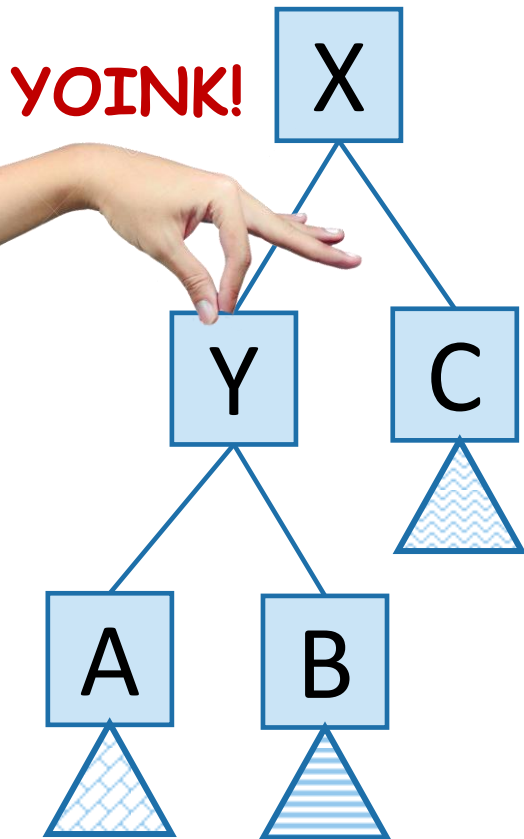
Self-Balancing Binary Search Trees



Idea 1: Rotations


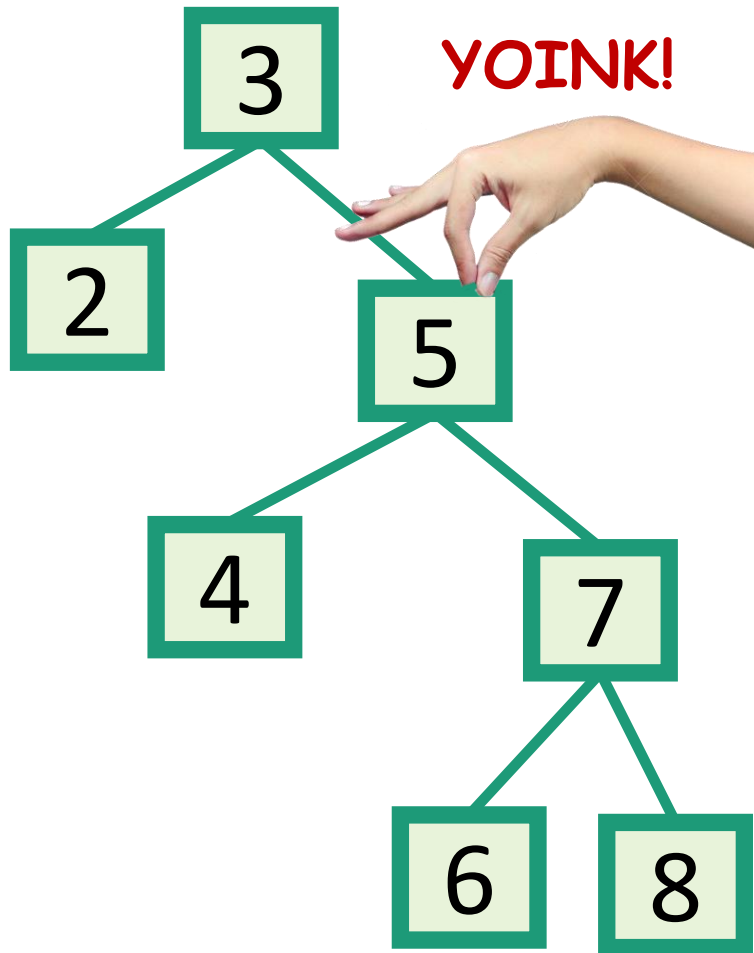
No matter what lives underneath A,B,C,
this takes time $O(1)$. (Why?)

- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

This seems helpful



Strategy?

- Whenever something seems unbalanced, do rotations until it's okay again.



This is pretty vague.

What do we mean by
“seems unbalanced”?

What’s “okay”?

Idea 2: have some proxy for balance

- Maintaining **perfect balance** is too hard.
- Instead, come up with some **proxy for balance**:
 - If the tree satisfies **[SOME PROPERTY]**, then it's pretty balanced.
 - We can maintain **[SOME PROPERTY]** using rotations.

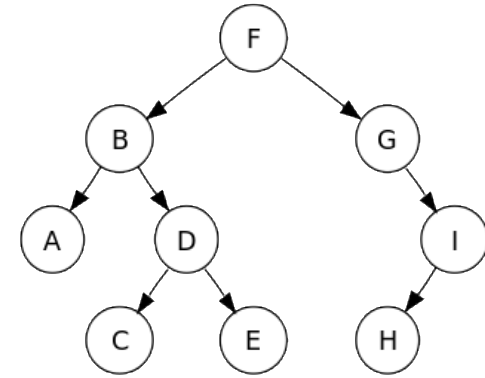


There are actually several ways to do this, but we'll see:

1. **AVL Tree** (Covered in DSA)
2. **Multiway-Search Tree (2-4 Tree)**
3. **Red-Black Tree**

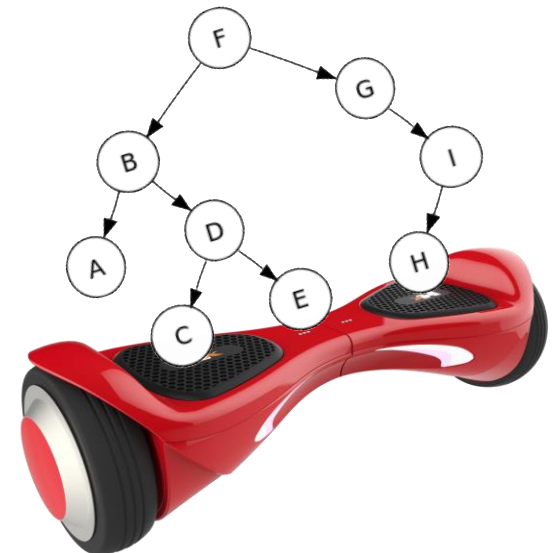
Today

- Begin a brief foray into data structures!
- Binary search trees
 - They are better when they're balanced.



this will lead us to...

- Self-Balancing Binary Search Trees
 - AVL Tree
 - Multiway-Search Tree
 - Red-Black Tree



Acknowledgement

- Stanford University