

Relational Database Design

Relational Database Design

- First Normal Form
- Pitfalls in Relational Database Design
- Functional Dependencies
- Decomposition
- Boyce-Codd Normal Form
- Third Normal Form
- Multivalued Dependencies and Fourth Normal Form
- Overall Database Design Process

First Normal Form

- Domain is **atomic** if its elements are considered to be indivisible units
 - Examples of non-atomic domains:
 - Set of names, composite attributes
 - Identification numbers like CS101 that can be broken up into parts
- A relational schema R is in **first normal form** if the domains of all attributes of R are atomic
- Non-atomic values complicate storage and encourage redundant (repeated) storage of data
 - E.g. Set of accounts stored with each customer, and set of owners stored with each account

First Normal Form (Contd.)

- Atomicity is actually a property of how the elements of the domain are used.

For example:

- Strings would normally be considered **indivisible**
- Suppose that students are given roll numbers which are strings of the form *CS0012* or *EE1127*
- If the first two characters are extracted to find the department, the domain of roll numbers is not atomic

Pitfalls in Relational Database Design

- Relational database design requires that we find a “good” collection of relation schemas. A bad design may lead to
 - Repetition of Information.
 - Inability to represent certain information.
- Design Goals:
 - Avoid redundant data
 - Ensure that relationships among attributes are represented
 - Facilitate the checking of updates for violation of database integrity constraints.

Example

- Consider the relation schema:

*lending-schema = (branch-name, branch-city, assets,
customer-name, loan-number, amount)*

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Downtown	Brooklyn	9000000	Jones	L-17	1000
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Downtown	Brooklyn	9000000	Jackson	L-14	1500

- Redundancy:**

- Data for *branch-name*, *branch-city*, *assets* are repeated for each loan
- Complicates updating, introducing possibility of inconsistency of *assets* value

- Null values**

- Cannot store information about a branch if no loans exist
- Can use null values, but they are difficult to handle.

Redundancy – Drawbacks

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Downtown	Brooklyn	9000000	Jones	L-17	1000
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Downtown	Brooklyn	9000000	Jackson	L-14	1500

- **Redundancy Storage:** Downtown repeated
?
- **Update Anomalies:** update assets where customer-name=Jones
?
- **Insertion Anomalies:** without loan-number cannot insert
?
- **Deletion Anomalies:** delete Smith record

Decomposition

- Decompose the relation schema *Lending-schema* into:

Branch-schema = (*branch-name*, *branch-city*, *assets*)

Loan-info-schema = (*customer-name*, *loan-number*, *branch-name*, *amount*)

- All attributes of an original schema (R) must appear in the decomposition (R_1, R_2):

$$R = R_1 \cup R_2$$

- Lossless-join decomposition:** instances

For all possible relations r on schema R

$$r = \Pi_{R_1}(r) \bowtie \Pi_{R_2}(r)$$

- Dependency-Preservation:** constraints

Example of Non Lossless-Join Decomposition

- Decomposition of $R = (A, B)$

$$R_1 = (A) \quad R_2 = (B)$$

A	B
α	1
α	2
β	1

r

A
α
β

$\Pi_A(r)$

B
1
2

$\Pi_B(r)$

$$\Pi_A(r) \bowtie \Pi_B(r)$$

A	B
α	1
α	2
β	1
β	2

Decomposition of $R = (A, B, C)$: $R_1 = (A, B)$ and $R_2 = (B, C)$

<i>A</i>	<i>B</i>	<i>C</i>

Goal — Devise a Theory for the Following

- Decide whether a particular relation R is in “good” form.
- In the case that a relation R is not in “good” form, decompose it into a set of relations $\{R_1, R_2, \dots, R_n\}$ such that
 - each relation is in good form
 - the decomposition is a lossless-join decomposition
- Our theory is based on:
 - functional dependencies
 - multivalued dependencies

Functional Dependencies

- Constraints on the set of legal relations.
- Require that the value for a certain set of attributes determines uniquely the value for another set of attributes.
- A functional dependency is a generalization of the notion of a *key*.

Functional Dependencies (Cont.)

- Let R be a relation schema $\alpha \subseteq R$ and $\beta \subseteq R$
- Functional dependency** $\alpha \rightarrow \beta$ holds on R **if and only if** for any legal relations $r(R)$, whenever any two tuples t_1 and t_2 of r agree on the attributes α , they also agree on the attributes β .
That is, $t_1[\alpha] = t_2[\alpha] \Rightarrow t_1[\beta] = t_2[\beta]$
- Example:** Consider $r(A,B)$ with the following instance of r .

A	B
1	4
1	5
3	7

- On this instance, $A \rightarrow B$ does **NOT** hold, but $B \rightarrow A$ does hold.

Functional Dependencies

- Example: $R = (A, B, C, D)$

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>

Functional Dependencies (Cont.)

- K is a **super-key** for a relation schema R if and only if $K \rightarrow R$
- K is a **candidate key** for R if and only if
 - $K \rightarrow R$, and
 - for no $\alpha \subset K$, $\alpha \rightarrow R$
- Functional dependencies allow us to express constraints that cannot be expressed using superkeys.

Consider the schema:

*Loan-info-schema = (customer-name, loan-number,
branch-name, amount)*

We expect this set of functional dependencies to hold:

loan-number \rightarrow amount

loan-number \rightarrow branch-name

but would not expect the following to hold:

loan-number \rightarrow customer-name

Use of Functional Dependencies

- We use functional dependencies to:
 - test relations to see if they are legal under a given set of functional dependencies.
 - If a relation r is legal under a set F of functional dependencies, we say that r **satisfies** F .
 - specify constraints on the set of legal relations
 - We say that F **holds on** R if all legal relations on R satisfy the set of functional dependencies F .
- **Note:** A specific instance of a relation schema may satisfy a functional dependency even if the functional dependency does not hold on all legal instances.

For example, a specific instance of *Loan-schema* may, by chance, satisfy
loan-number \rightarrow *customer-name*

Functional Dependencies (Cont.)

- A functional dependency is **trivial** if it is satisfied by all instances of a relation
 - *E.g.*
 - *customer-name, loan-number \rightarrow customer-name*
 - *customer-name \rightarrow customer-name*
 - In general, $\alpha \rightarrow \beta$ is trivial if $\beta \subseteq \alpha$

BREAK

Closure of a Set of Functional Dependencies

- Given a set F , set of functional dependencies, there are certain other functional dependencies that are logically implied by F .
 - E.g. If $A \rightarrow B$ and $B \rightarrow C$, then we can infer that $A \rightarrow C$
- Set of all functional dependencies logically implied by F is the *closure* of F .
- We denote the *closure* of F by F^+ .
- We can find all of F^+ by applying Armstrong's Axioms:
 - if $\beta \subseteq \alpha$, then $\alpha \rightarrow \beta$ (**reflexivity**)
 - if $\alpha \rightarrow \beta$, then $\gamma \alpha \rightarrow \gamma \beta$ (**augmentation**)
 - if $\alpha \rightarrow \beta$, and $\beta \rightarrow \gamma$, then $\alpha \rightarrow \gamma$ (**transitivity**)
- These rules are
 - sound** (generate only functional dependencies that actually hold) and
 - complete** (generate all functional dependencies that hold).

Example

- $R = (A, B, C, G, H, I)$

$$F = \left\{ \begin{array}{l} A \rightarrow B \\ A \rightarrow C \\ CG \rightarrow H \\ CG \rightarrow I \\ B \rightarrow H \end{array} \right\}$$

- **some members of F^+**

- $A \rightarrow H$

- by transitivity from $A \rightarrow B$ and $B \rightarrow H$

- $AG \rightarrow I$

- by augmenting $A \rightarrow C$ with G , to get $AG \rightarrow CG$
and then transitivity with $CG \rightarrow I$

- $CG \rightarrow HI$

- from $CG \rightarrow H$ and $CG \rightarrow I$: “union rule” can be inferred from
 - definition of functional dependencies, or
 - Augmentation of $CG \rightarrow I$ to infer $CG \rightarrow CGI$, augmentation of $CG \rightarrow H$ to infer $CGI \rightarrow HI$, and then transitivity

Procedure for Computing F^+

- To compute the closure of a set of functional dependencies F :
- $F^+ = F$
repeat
 for each functional dependency f in F^+
 apply reflexivity and augmentation rules on f
 add the resulting functional dependencies to F^+
 for each pair of functional dependencies f_1 and f_2 in F^+
 if f_1 and f_2 can be combined using transitivity
 then add the resulting functional dependency to F^+
until F^+ does not change any further

NOTE: We will see an alternative procedure for this task later

Closure of Functional Dependencies (Cont.)

- We can further simplify manual computation of F^+ by using the following additional rules.
 - If $\alpha \rightarrow \beta$ holds and $\alpha \rightarrow \gamma$ holds, then $\alpha \rightarrow \beta\gamma$ holds
(**union**)
 - If $\alpha \rightarrow \beta\gamma$ holds, then $\alpha \rightarrow \beta$ holds and $\alpha \rightarrow \gamma$ holds
(**decomposition**)
 - If $\alpha \rightarrow \beta$ holds and $\gamma\beta \rightarrow \delta$ holds, then $\alpha\gamma \rightarrow \delta$ holds
(**pseudotransitivity**)

The above rules can be inferred from Armstrong's axioms.

Closure of Attribute Sets

- Given a set of attributes α , define the *closure* of α under F (denoted by α^+) as the set of attributes that are functionally determined by α under F :

$$\alpha \rightarrow \beta \text{ is in } F^+ \iff \beta \subseteq \alpha^+$$

- Algorithm to compute α^+ , the closure of α under F

result := α ;

while (changes to *result*) **do**

for each $\beta \rightarrow \gamma$ **in** F **do**

begin

if $\beta \subseteq \text{result}$ **then** *result* := *result* $\cup \gamma$

end

Example of Attribute Set Closure

- $R = (A, B, C, G, H, I)$

- $F = \{ \begin{array}{l} A \rightarrow B \\ A \rightarrow C \\ CG \rightarrow H \\ CG \rightarrow I \\ B \rightarrow H \end{array} \}$

- $(AG)^+$

1. $result = AG$

2. $result = ABCG$ ($A \rightarrow C$ and $A \rightarrow B$)

3. $result = ABCGH$ ($CG \rightarrow H$ and $CG \subseteq AGBC$)

4. $result = ABCGHI$ ($CG \rightarrow I$ and $CG \subseteq AGBCH$)

- Is AG a candidate key?

1. Is AG a super key?

1. Does $AG \rightarrow R$?

2. Is any subset of AG a superkey?

1. Does $A^+ \rightarrow R$?

2. Does $G^+ \rightarrow R$?

Uses of Attribute Closure

There are several uses of the attribute closure algorithm:

- **Testing for superkey:**
 - To test if α is a superkey, we compute α^+ , and check if α^+ contains all attributes of R .
- **Testing functional dependencies**
 - To check if a functional dependency $\alpha \rightarrow \beta$ holds (or, in other words, is in F^+), just check if $\beta \subseteq \alpha^+$.
 - That is, we compute α^+ by using attribute closure, and then check if it contains β .
 - Is a simple and cheap test, and very useful
- **Computing closure of F**
 - For each $\gamma \subseteq R$, we find the closure γ^+ , and for each $S \subseteq \gamma^+$, we output a functional dependency $\gamma \rightarrow S$.

THANK YOU

Reference (Textbook):

- Silberschatz, H. Korth & S. Sudarshan, Database System Concepts, McGraw-Hill Education, 6th Edition, 2010